

\$ COST

**Cost cutting and
efficiency are not the
same thing.**

Most teams optimize the bill.
The bill is the symptom.

EFFICIENCY

Cutting cost shrinks the system.
Efficiency shrinks the waste.

COST CUT




system shrinks

Throttle the model, cap the context, ship worse output – costs drop. So does the product.


Efficiency removes work the system never needed to do.

EFFICIENCY




only waste removed

**ScreenWeave: 91% cost reduction,
zero capability removed.**



We didn't downgrade the model or starve the context.
We cut the redundant work the agent kept repeating.

Where the waste actually lives.



Re-sending context. Re-running steps.
Paying premium rates for floor-level work.

**Efficiency is an architecture
decision, not a settings toggle.**



You don't find 91% in a config flag.
You find it in how context, routing, and state are designed.

The test: cut the bill, does the output survive?



If quality drops, you cut cost.
If quality holds, you cut waste.
Only one of those scales.

Measure waste, not spend.

Spend tells you what you paid.
Waste tells you what you can remove
without losing anything.

**What's the most expensive redundant step in
your agent pipeline right now – and is anyone
measuring it?**

AIWeave / applied research – aiweave.org
